

Underwater Image Enhancement - Shreyas Lal

Subtask 1: Variational Autoencoders (VAEs)

Import necessary libraries

```
import os
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import load_img,
img_to_array
import tensorflow as tf
from skimage.metrics import structural_similarity as ssim
```

Set paths

```
train_raw_path = '/kaggle/input/slal-task2-dataset/Train/Raw/'
train_ref_path = '/kaggle/input/slal-task2-dataset/Train/Reference/'
```

Load Images

Function to load and preprocess images

```
def load_images(path, target_size=(480, 640)):
    images = [] # Initialize an empty list to store images
    for file in os.listdir(path): # Loop through each file in the
        directory
            # Ensure the file is a valid image
            if file.endswith(('.png', '.jpg', '.jpeg')): # Only load
                image files
                    img = load_img(os.path.join(path, file), target_size=(480,
640)) # Load and resize image
                    img = img_to_array(img) / 255.0 # Convert to array and
normalize pixel values to [0, 1]
                    images.append(img) # Add the preprocessed image to the
list
    return np.array(images) # Convert list to numpy array and return
```

Load training images and reference images

```
train_images = load_images(train_raw_path)
reference_images = load_images(train_ref_path)
```

Visualize a few loaded images

```
for i in range(min(5, len(train_images))):  
    plt.imshow(train_images[i])  
    plt.axis('off')  
    plt.title(f'Training Image {i+1}')  
    plt.show()
```

Training Image 1



Training Image 2



Training Image 3



Training Image 4



Training Image 5




```

# Define the VAE model
def build_vae():
    inputs = layers.Input(shape=(480, 640, 3)) # Input shape for
images

    # Encoder
    x = layers.Conv2D(32, 3, activation='relu', padding='same')
(inputs) # 1st convolutional layer
    x = layers.MaxPooling2D((2, 2))(x) # Downsample the feature maps

    x = layers.Conv2D(64, 3, activation='relu', padding='same')(x) #
2nd convolutional layer
    x = layers.MaxPooling2D((2, 2))(x) # Further downsampling

    x = layers.Conv2D(128, 3, activation='relu', padding='same')(x) #
3rd convolutional layer
    x = layers.MaxPooling2D((2, 2))(x) # Further downsampling

    # Latent space
    x = layers.Flatten()(x) # Flatten the output
    z_mean = layers.Dense(64)(x) # Mean of the latent space
    z_log_var = layers.Dense(64)(x) # Log variance of the latent
space

    # Sampling layer
    z = layers.Lambda(sampling, output_shape=(64,))([z_mean,
z_log_var]) # Sample from the latent space

    # Decoder
    x = layers.Dense(30 * 40 * 128, activation='relu')(z) # Expand
the latent representation
    x = layers.Reshape((30, 40, 128))(x) # Reshape to a 3D tensor

    # Deconvolutional layers to upsample back to the original image
size
    x = layers.Conv2DTranspose(128, 3, activation='relu',
padding='same')(x) # 1st upsampling step
    x = layers.UpSampling2D((2, 2))(x) # Upsample the feature maps

    x = layers.Conv2DTranspose(64, 3, activation='relu',
padding='same')(x) # 2nd upsampling step
    x = layers.UpSampling2D((2, 2))(x) # Upsample again

    x = layers.Conv2DTranspose(32, 3, activation='relu',
padding='same')(x) # 3rd upsampling step
    x = layers.UpSampling2D((2, 2))(x) # Upsample again

    # Final layer to output the reconstructed image
    outputs = layers.Conv2DTranspose(3, 3, activation='sigmoid',
padding='same')(x) # 3 channels for RGB

```

```

vae = models.Model(inputs, outputs) # Create the VAE model
return vae

# Function for sampling from the latent space
def sampling(args):
    z_mean, z_log_var = args # Get mean and log variance
    batch = tf.shape(z_mean)[0] # Get the batch size
    dim = tf.shape(z_mean)[1] # Get the dimension of the latent space
    epsilon = tf.keras.backend.random_normal(shape=(batch, dim)) #
    Sample epsilon
    return z_mean + tf.exp(0.5 * z_log_var) * epsilon # Return the
    sampled latent vector

# Load images
train_raw_path = '/kaggle/input/slal-task2-dataset/Train/Raw/'
train_images = load_images(train_raw_path)

# Compile and train the VAE
vae = build_vae() # Build the VAE model
vae.compile(optimizer='adam', loss='binary_crossentropy') # Compile
model with Adam optimizer and binary crossentropy loss
vae.fit(train_images, train_images, epochs=50, batch_size=32) # Train
the VAE

# Function to generate enhanced images from the VAE
def generate_images(model, n=10):
    noise = np.random.normal(size=(n, 64)) # Generate noise in the
    latent space size (64)
    generated_images = model.predict(noise) # Use the VAE model to
    generate images from the noise
    return generated_images

# Generate and visualize the enhanced images
generated_images = generate_images(vae, n=5)

# Evaluate and compare original and generated images
def evaluate_images(original, generated):
    mse_scores = [] # Store Mean Squared Error scores
    ssim_scores = [] # Store Structural Similarity Index scores
    for orig, gen in zip(original, generated):
        mse = mean_squared_error(orig.flatten(), gen.flatten()) #
        Calculate MSE
        ssim_value = ssim(orig, gen, multichannel=True) # Calculate
        SSIM
        mse_scores.append(mse) # Append MSE score
        ssim_scores.append(ssim_value) # Append SSIM score
    return mse_scores, ssim_scores

# Evaluate the first few original images against generated images

```

```
mse_scores, ssim_scores = evaluate_images(train_images[:5],
generated_images)
```

```
# Print evaluation results
```

```
print("Mean Squared Errors (MSE):", mse_scores)
```

```
print("Structural Similarity Indices (SSIM):", ssim_scores)
```

```
# Visualize original vs generated images
```

```
def visualize_comparison(original, generated):
```

```
    n = min(len(original), len(generated))
```

```
    plt.figure(figsize=(15, 6))
```

```
    for i in range(n):
```

```
        plt.subplot(2, n, i + 1)
```

```
        plt.imshow(original[i]) # Original image
```

```
        plt.axis('off') # Turn off axis labels
```

```
        plt.title("Original")
```

```
        plt.subplot(2, n, i + 1 + n)
```

```
        plt.imshow(generated[i]) # Generated image
```

```
        plt.axis('off') # Turn off axis labels
```

```
        plt.title("Generated")
```

```
    plt.show()
```

```
# Compare the first few original and generated images
```

```
visualize_comparison(train_images[:5], generated_images)
```

Epoch 1/50

```
-----
-----
ValueError                                Traceback (most recent call
last)
Cell In[12], line 58
     56 vae = build_vae() # Build the VAE model
     57 vae.compile(optimizer='adam', loss='binary_crossentropy') #
Compile model with Adam optimizer and binary crossentropy loss
--> 58 vae.fit(train_images, train_images, epochs=50, batch_size=32)
# Train the VAE
     60 # Function to generate enhanced images from the VAE
     61 def generate_images(model, n=10):

File
/opt/conda/lib/python3.10/site-packages/keras/src/utils/traceback_util
s.py:122, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    119     filtered_tb = _process_traceback_frames(e.__traceback__)
    120     # To get the full stack trace, call:
    121     # `keras.config.disable_traceback_filtering()`
--> 122     raise e.with_traceback(filtered_tb) from None
    123 finally:
    124     del filtered_tb
```

```

File
/opt/conda/lib/python3.10/site-packages/keras/src/backend/tensorflow/
nn.py:668, in binary_crossentropy(target, output, from_logits)
    666 for e1, e2 in zip(target.shape, output.shape):
    667     if e1 is not None and e2 is not None and e1 != e2:
--> 668         raise ValueError(
    669             "Arguments `target` and `output` must have the
same shape. "
    670             "Received: "
    671             f"target.shape={target.shape},
output.shape={output.shape}"
    672         )
    674 output, from_logits = _get_logits(
    675     output, from_logits, "Sigmoid", "binary_crossentropy"
    676 )
    678 if from_logits:

```

```

ValueError: Arguments `target` and `output` must have the same shape.
Received: target.shape=(None, 480, 640, 3), output.shape=(None, 240,
320, 3)

```

VAE Modelling

```

# Sampling function for the reparameterization trick in VAEs
def sampling(args):
    z_mean, z_log_var = args
    batch = tf.shape(z_mean)[0]
    dim = tf.shape(z_mean)[1]
    epsilon = tf.keras.backend.random_normal(shape=(batch, dim))
    return z_mean + tf.exp(0.5 * z_log_var) * epsilon

# Build the VAE model
def build_vae():
    # Input shape adjusted for 640x480 RGB images
    inputs = layers.Input(shape=(480, 640, 3))

    # Encoder
    x = layers.Conv2D(32, 3, activation='relu', padding='same')(
inputs)
    x = layers.MaxPooling2D()(x)
    x = layers.Conv2D(64, 3, activation='relu', padding='same')(x)
    x = layers.MaxPooling2D()(x)

    # Latent space representation
    x = layers.Flatten()(x)
    x = layers.Dense(128, activation='relu')(x) # Increased dense
layer for complexity

```



```

# Mean and log variance for the latent space
z_mean = layers.Dense(32)(x)
z_log_var = layers.Dense(32)(x)

# Sample a point in latent space
z = layers.Lambda(sampling)([z_mean, z_log_var])

# Decoder
x = layers.Dense(120 * 160 * 64, activation='relu')(z) # Adjust
based on downsampling
x = layers.Reshape((120, 160, 64))(x)

x = layers.Conv2DTranspose(64, 3, activation='relu',
padding='same')(x)
x = layers.UpSampling2D()(x)
x = layers.Conv2DTranspose(32, 3, activation='relu',
padding='same')(x)
x = layers.UpSampling2D()(x)

# Final layer to output the reconstructed image with the same
shape as the input
outputs = layers.Conv2DTranspose(3, 3, activation='sigmoid',
padding='same')(x)

# Define the VAE model
vae = models.Model(inputs, outputs)

# Encoder and decoder models for optional use
encoder = models.Model(inputs, z_mean)

# Decoder input is latent space, then generate the image
latent_inputs = layers.Input(shape=(32,))
x = layers.Dense(120 * 160 * 64, activation='relu')(latent_inputs)
x = layers.Reshape((120, 160, 64))(x)
x = layers.Conv2DTranspose(64, 3, activation='relu',
padding='same')(x)
x = layers.UpSampling2D()(x)
x = layers.Conv2DTranspose(32, 3, activation='relu',
padding='same')(x)
x = layers.UpSampling2D()(x)
decoded_output = layers.Conv2DTranspose(3, 3,
activation='sigmoid', padding='same')(x)
decoder = models.Model(latent_inputs, decoded_output)

return vae, encoder, decoder

# Compile and train the VAE
vae, encoder, decoder = build_vae()

```

```

vae.compile(optimizer='adam', loss='binary_crossentropy')

# Assuming 'train_images' is your dataset of underwater images
# Normalize train_images to the range [0, 1]
train_images = train_images / 255.0

vae.fit(train_images, train_images, epochs=50, batch_size=32)

# Function to generate images from random latent space vectors
def generate_images(decoder, n=10):
    noise = np.random.normal(size=(n, 32)) # Generate random vectors
    in latent space
    generated_images = decoder.predict(noise) # Generate images using
    decoder
    return generated_images

# Generate and visualize the generated images
generated_images = generate_images(decoder)

for i in range(len(generated_images)):
    plt.imshow(generated_images[i])
    plt.axis('off')
    plt.show()

def build_vae():
    # Define the input shape for the encoder; here it's set for
    128x128 RGB images
    inputs = layers.Input(shape=(128, 128, 3))

    # Encoder
    #Start with a convolutional layer followed by max pooling
    x = layers.Conv2D(32, 3, activation='relu', padding='same')
    (inputs)# 32 filters, kernel size of 3x3
    x = layers.MaxPooling2D()(x) # Downsample the feature maps

    # Another convolutional layer to extract more features
    x = layers.Conv2D(64, 3, activation='relu', padding='same')(x) #
    64 filters
    x = layers.MaxPooling2D()(x) # Further downsampling

    # Latent space
    #Flatten the output from the convolutional layers to create a
    vector representation
    latent_space = layers.Flatten()(x) # Convert to a 1D vector
    latent_space = layers.Dense(32)(latent_space) # Dense layer for
    latent representation

    # Decoder
    # Reconstruct the image from the latent space
    x = layers.Dense(64 * 32 * 32, activation='relu')(latent_space) #

```

```

Dense layer to expand the latent representation
x = layers.Reshape((32, 32, 64))(x) # Reshape into a 3D tensor
(32x32 feature maps with 64 channels)

# Deconvolutional layers to upsample back to the original image
size
x = layers.Conv2DTranspose(64, 3, activation='relu',
padding='same')(x) # First upsampling step
x = layers.UpSampling2D()(x) # Upsample the feature maps

x = layers.Conv2DTranspose(32, 3, activation='relu',
padding='same')(x) # Second upsampling step
x = layers.UpSampling2D()(x) # Unsample again

# Final layer to output the reconstructed image with same
dimensions as input
outputs = layers.Conv2DTranspose(3, 3, activation='sigmoid',
padding='same')(x)

# Create the VAE model combining inputs and outputs
vae = models.Model(inputs, outputs)
return vae

# Compile and train the VAE
vae = build_vae() # Build the VAE model
vae.compile(optimizer='adam', loss='binary_crossentropy') # Compile
model with Adam optimizer and binary crossentropy loss
vae.fit(train_images, train_images, epochs=50, batch_size=32) # Train
the VAE, using images as both input and output

# Function to generate enhanced images from the VAE
def generate_images(model, n=10):
    # Create an array of random noise with the same shape as the
    images
    noise = np.random.normal(size=(n, 128, 128, 3)) # Sample random
    noise
    generated_images = model.predict(noise) # Use the VAE model to
    generate images from the noise
    return generated_images

# Generate and visualize images
generated_images = generate_images(vae) # Call the function to
generate a specified number of images
for i in range(len(generated_images)): # Loop through generated images
    plt.imshow(generated_images[i])
    plt.axis('off')
    plt.show()

Epoch 1/50
22/22 _____ 25s 985ms/step - loss: 0.6609

```

| | | | |
|-------------|----------------------|----------------|--------------|
| Epoch 2/50 | | | |
| 22/22 | ████████████████████ | 21s 931ms/step | loss: 0.6310 |
| Epoch 3/50 | | | |
| 22/22 | ████████████████████ | 42s 968ms/step | loss: 0.6272 |
| Epoch 4/50 | | | |
| 22/22 | ████████████████████ | 41s 952ms/step | loss: 0.6244 |
| Epoch 5/50 | | | |
| 22/22 | ████████████████████ | 41s 959ms/step | loss: 0.6160 |
| Epoch 6/50 | | | |
| 22/22 | ████████████████████ | 21s 935ms/step | loss: 0.6054 |
| Epoch 7/50 | | | |
| 22/22 | ████████████████████ | 21s 935ms/step | loss: 0.5972 |
| Epoch 8/50 | | | |
| 22/22 | ████████████████████ | 41s 949ms/step | loss: 0.5949 |
| Epoch 9/50 | | | |
| 22/22 | ████████████████████ | 21s 932ms/step | loss: 0.5884 |
| Epoch 10/50 | | | |
| 22/22 | ████████████████████ | 21s 970ms/step | loss: 0.5885 |
| Epoch 11/50 | | | |
| 22/22 | ████████████████████ | 22s 1s/step | loss: 0.5834 |
| Epoch 12/50 | | | |
| 22/22 | ████████████████████ | 40s 955ms/step | loss: 0.5843 |
| Epoch 13/50 | | | |
| 22/22 | ████████████████████ | 20s 925ms/step | loss: 0.5786 |
| Epoch 14/50 | | | |
| 22/22 | ████████████████████ | 21s 967ms/step | loss: 0.5787 |
| Epoch 15/50 | | | |
| 22/22 | ████████████████████ | 40s 923ms/step | loss: 0.5820 |
| Epoch 16/50 | | | |
| 22/22 | ████████████████████ | 21s 957ms/step | loss: 0.5718 |
| Epoch 17/50 | | | |
| 22/22 | ████████████████████ | 20s 929ms/step | loss: 0.5743 |
| Epoch 18/50 | | | |
| 22/22 | ████████████████████ | 41s 942ms/step | loss: 0.5757 |
| Epoch 19/50 | | | |
| 22/22 | ████████████████████ | 41s 946ms/step | loss: 0.5732 |
| Epoch 20/50 | | | |
| 22/22 | ████████████████████ | 21s 943ms/step | loss: 0.5741 |
| Epoch 21/50 | | | |
| 22/22 | ████████████████████ | 41s 941ms/step | loss: 0.5743 |
| Epoch 22/50 | | | |
| 22/22 | ████████████████████ | 21s 936ms/step | loss: 0.5770 |
| Epoch 23/50 | | | |
| 22/22 | ████████████████████ | 41s 954ms/step | loss: 0.5741 |
| Epoch 24/50 | | | |
| 22/22 | ████████████████████ | 20s 928ms/step | loss: 0.5635 |
| Epoch 25/50 | | | |
| 22/22 | ████████████████████ | 20s 907ms/step | loss: 0.5668 |
| Epoch 26/50 | | | |

| | | | | | |
|-------------|-------|-----|------------|---|--------------|
| 22/22 | _____ | 21s | 935ms/step | - | loss: 0.5685 |
| Epoch 27/50 | | | | | |
| 22/22 | _____ | 20s | 902ms/step | - | loss: 0.5660 |
| Epoch 28/50 | | | | | |
| 22/22 | _____ | 21s | 949ms/step | - | loss: 0.5707 |
| Epoch 29/50 | | | | | |
| 22/22 | _____ | 40s | 910ms/step | - | loss: 0.5612 |
| Epoch 30/50 | | | | | |
| 22/22 | _____ | 21s | 945ms/step | - | loss: 0.5674 |
| Epoch 31/50 | | | | | |
| 22/22 | _____ | 21s | 955ms/step | - | loss: 0.5678 |
| Epoch 32/50 | | | | | |
| 22/22 | _____ | 41s | 947ms/step | - | loss: 0.5663 |
| Epoch 33/50 | | | | | |
| 22/22 | _____ | 41s | 935ms/step | - | loss: 0.5666 |
| Epoch 34/50 | | | | | |
| 22/22 | _____ | 41s | 921ms/step | - | loss: 0.5629 |
| Epoch 35/50 | | | | | |
| 22/22 | _____ | 21s | 946ms/step | - | loss: 0.5572 |
| Epoch 36/50 | | | | | |
| 22/22 | _____ | 21s | 955ms/step | - | loss: 0.5650 |
| Epoch 37/50 | | | | | |
| 22/22 | _____ | 21s | 949ms/step | - | loss: 0.5592 |
| Epoch 38/50 | | | | | |
| 22/22 | _____ | 41s | 941ms/step | - | loss: 0.5591 |
| Epoch 39/50 | | | | | |
| 22/22 | _____ | 41s | 945ms/step | - | loss: 0.5600 |
| Epoch 40/50 | | | | | |
| 22/22 | _____ | 20s | 923ms/step | - | loss: 0.5650 |
| Epoch 41/50 | | | | | |
| 22/22 | _____ | 21s | 962ms/step | - | loss: 0.5632 |
| Epoch 42/50 | | | | | |
| 22/22 | _____ | 21s | 952ms/step | - | loss: 0.5619 |
| Epoch 43/50 | | | | | |
| 22/22 | _____ | 20s | 915ms/step | - | loss: 0.5648 |
| Epoch 44/50 | | | | | |
| 22/22 | _____ | 21s | 958ms/step | - | loss: 0.5619 |
| Epoch 45/50 | | | | | |
| 22/22 | _____ | 41s | 936ms/step | - | loss: 0.5627 |
| Epoch 46/50 | | | | | |
| 22/22 | _____ | 22s | 996ms/step | - | loss: 0.5614 |
| Epoch 47/50 | | | | | |
| 22/22 | _____ | 41s | 971ms/step | - | loss: 0.5632 |
| Epoch 48/50 | | | | | |
| 22/22 | _____ | 22s | 1s/step | - | loss: 0.5602 |
| Epoch 49/50 | | | | | |
| 22/22 | _____ | 21s | 937ms/step | - | loss: 0.5589 |
| Epoch 50/50 | | | | | |

22/22 ————— 42s 963ms/step - loss: 0.5597
1/1 ————— 0s 313ms/step









